

Neurokernel Core Architecture

Current Design, Limitations, and Future Aims

Lev Givon

Bionet Group
Department of Electrical Engineering
Columbia University

March 17, 2016



Outline

- 1 Interfacing LPU's
- 2 Communication
- 3 Neurodriver
- 4 Future Aims

Port Identifiers and Selectors

- Each port must be assigned a unique *identifier*.
- Mandatory path-like format: `/level10/level11/...`
- *Selectors* represent multiple identifiers: `/level10[0:5]`
- Syntax goodies: ranges, lists, concatenation, outer products, elementwise products.
- **Selector** class adds support for set operations.

`nk.plsel`

- **Selector** - validates/expands/caches selectors
 - **SelectorParser** - selector grammar
 - **SelectorMethods** - selector manipulation routines
-
- Selectors expanded by parser into lists of tuples. Expensive!
 - Adding ports to existing interface also expensive - avoid by creating with all required ports.

Selector Syntax Features

Identifier/Selector	Comments
<code>/med/L1[0]</code>	selects a single port
<code>/med/L1/0</code>	equivalent to <code>/med/L1[0]</code>
<code>/med+/L1[0]</code>	equivalent to <code>/med/L1[0]</code>
<code>/med/[L1, L2][0]</code>	selects two ports
<code>/med/L1[0, 1]</code>	another example of two ports
<code>/med/L1[0], /med/L1[1]</code>	equivalent to <code>/med/L1[0, 1]</code>
<code>/med/L1[0:10]</code>	selects ten ports
<code>/med/L1/*</code>	selects all ports starting with <code>/med/L1</code>
<code>(/med/L1, /med/L2) + [0]</code>	equivalent to <code>/med/[L1, L2][0]</code>
<code>/med/[L1, L2].+[0:2]</code>	equivalent to <code>/med/L1[0], /med/L2[1]</code>

LPU Interface Design

- *Interface* = port identifiers + port attributes + port \leftrightarrow transmitted data array mapping.
- **Interface** class can encapsulate multiple interfaces.
- Required attributes: interface identifier, I/O direction (input or output), type (spiking or graded potential).
- Stored in Pandas **DataFrame**: attribs \rightarrow data, (expanded) ports \rightarrow index.
- Interface compatibility: ports in both interfaces must have same type and inverse I/O attributes.

`nk.pattern`

- **Interface** - interface class

- Compatibility checking is expensive (requires equiv. of inner join)!

Port Maps

- Goal: use selectors to access array of data transmitted to/from ports.
- Solution: *map* identifiers to array indices: `/label10[a,b,c] → [0,1,2]`
- Stored in Pandas **DataFrame**: array indices → data, ports → frame index.
- Port data array must have atomic dtype (`int32`, `float64`, etc.).

`nk.pm`, `nk.pm_gpu`

- **PortMapper** - maps ports to host memory array
 - **GPUPortMapper** - maps ports to GPU memory array
-
- Selector-based access is expensive! Solution: use indices rather than selectors during model execution.

Inter-LPU Connectivity Patterns

- Each **Pattern** connects two LPU's - N patterns required to connect one LPU to N other LPU's.
- Each pattern contains two interfaces (stored in one **Interface** instance).
- Stored in Pandas **DataFrame**: index contains connected source/destination ports (expanded).
- If source/dest pair \neq pattern, they are not connected.
- Selectors may be used to set/select connections, but as before..

`nk.pattern`

- **Pattern** - pattern class

- .. selector-based access/adding new connections are expensive!
- Use special classmethods to create patterns from selectors.

More Areas for Improvement

- Selector expansion is done to enable storage in Pandas **MultiIndex**.
- The good: can facilitate selection of groups of ports using [hierarchical indexing](#).
- The bad: handling of multilevel identifiers is complicated and partially broken (see issue [#52](#)), expensive selection operations.
- Possibility: store identifiers as Pandas **Index** of string labels?
- Possibility: ensure that parsing/expansion never occurs during model execution?

Defining LPUs

- LPU models are implemented as Python classes with an **Interface** attribute and **run_step()** method.
- **run_step()** is invoked at each execution step.
- Incoming/outgoing port data must be read/updated from/to the interface port mapper within **run_step()**.
- After invocation of **run_step()**, Neurokernel *synchronizes* the LPU's ports with those of other connected LPUs.

nk.core, nk.core_gpu

- **Module** - parent LPU class
- GPU and non-GPU parent classes are separate - could be conflated.

Multi-LPU Emulation Setup

- LPU classes + params and patterns must be added to an emulation by a **Manager** class instance.
- NK uses MPI-2 *dynamic process creation* to automatically spawn sufficient # of processes to run emulation.
- Processes started/stopped by control messages via interprocess communicator.
- Run loop timing info also collected by manager via messages.

`nk.core, nk.core_gpu`

- **Manager** - emulation manager

More Areas for Improvement

- Process spawning requires active MPI env.
- Solution: relaunch main script via `mpi_relaunch` module and then spawn.
- Downside: can't develop/debug interactively.
- Possibility: decouple manager process from spawning process, but let former steer latter (cf. [IPython Parallel](#)).
- Objects required by LPU class need to be present in spawned proc namespace.
- Solution: recursively find/[serialize](#) params/globals accessed by class, transmit to spawned proc.
- Downside: some objects can't be serialized, transmission of large/many params/globals is time-consuming.

Inter-LPU Communication

- Communication between LPUs automatically performed by NK after each execution step:
 - 1 Source: output port data → transmission buffer.
 - 2 Source: buffer → MPI.
 - 3 Destination: MPI → buffer.
 - 4 Destination: buffer → input port data.
 - Data copied to/from contiguous buffers to enable CUDA-enabled MPI to use GPUDirect.
 - Transmissions are launched asynchronously, but each LPU waits until initiated send operations complete before next exec step.
-
- Noncontiguous copy to/from transmission buffers is inefficient.
 - Possibility: use multiple sends per LPU?
 - Possibility: fully asynchronous transmission (with deadlock breaking algorithm)?

Neurodriver I

- Neurodriver: configurable LPU implementation.
- Supports circuits comprising several point neuron models and associated synapse models.
- New models may be defined by subclassing **BaseNeuron**, **BaseSynapse**.
- Class may be instantiated with model parameters:

```
{'ModelName': {'Param0': [val0, val1, ..],
               'Param1': [val4, val5, ..]}, ...}}
```


- Can specify circuit as graph (GEXF, NetworkX).
- # of neurons and synapses only constrained by memory.
- All spiking model port data internally stored consecutively in one array, as is graded potential model port data.

Neurodriver II


`nk.LPU.LPU`, `nk.LPU.neurons`, `nk.LPU.synapses`

- **BaseNeuron**, **BaseSynapse**, etc. - neuron/synapse model implementations.
 - **LPU** - subclass of **Module**
-
- Legacy design: input ports must be explicitly specified, output ports specified implicitly by setting **public** model param.
 - Legacy design: synapses must be explicitly marked as belonging to different classes (spiking → graded potential, graded potential → spiking, etc.)
 - Extra noncontiguous GPU memory copies to/from port mappers.
 - Hard to add models with complex input/output relationships.
 - Possibility: make neurons/synapses subclasses of a more general *component* class (#2).

Realizing Vertical APIs



Realizing Vertical APIs



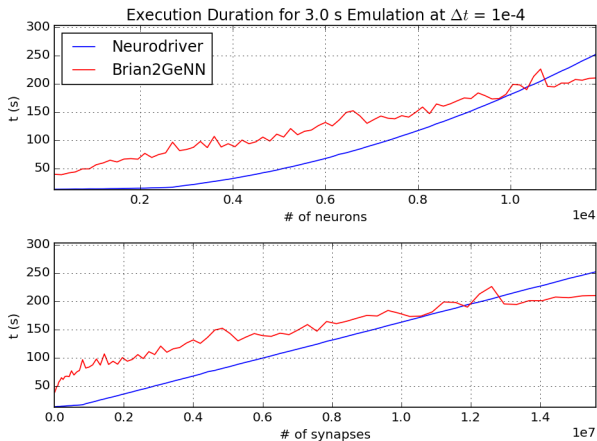
Improving GPU Resource Utilization

- Each LPU knows about one GPU (but multiple LPUs can use the same GPU).
 - Direct access to GPUs by LPU implementation precludes efficient resource use.
-
- Restrict direct access to GPUs to compute plane.
 - Add mechanism for mapping circuit components to resources (simple prototype using **METIS** already used in benchmarks).
 - Devise resource alloc policies that optimize over available GPUs, bus bandwidth, model component cost, etc.
 - Utilize structural data in NeuroArch for resource allocation.
 - Possibility: parameterize generation of CUDA code based model resources reqs.

Improving Model Component Support

- Neurodriver plugin system sufficient for point neurons/simple synapses, but hard to extend to more complex models.
 - Plugin system complicates efficient resource utilization.
-
- Generalize to support models with more complex I/O.
 - Only the application plane must know what an LPU is.
 - Add new primitives/components to compute plane - only exposed through vertical API.
 - Convert Neurodriver to compute plane engine.
 - Possibility: enforce declarative definition of LPUs?

Improving Emulation Performance



Hacking Notes

- Control logging with `nk.tools.logging.setup_logger`.
 - Logs may be written to screen, file, or both.
 - Logging can be turned off to prevent I/O slowdown.
 - Debugging spawned Python processes - see [ripdb](#).
 - Tip: when printing `var` contents to log,
 - 1 set `multiline=True` in `setup_logger()` params;
 - 2 make sure that `repr(var)` returns something sane!
 - If `debug=False` in `Module` constructor, exceptions in `run_step` are *not* fatal!
-
- Logged messages can get lost/overwritten due to MPI-IO sync issues.
 - Crude solution: pipe stdout to file.
 - Possibility: send logs to aggregator process rather than disk to avoid I/O slow-down.